

UNIX Networking 2


Benjamin Brewster

Adapted from Jon Herlocker, OSU

Network PIs

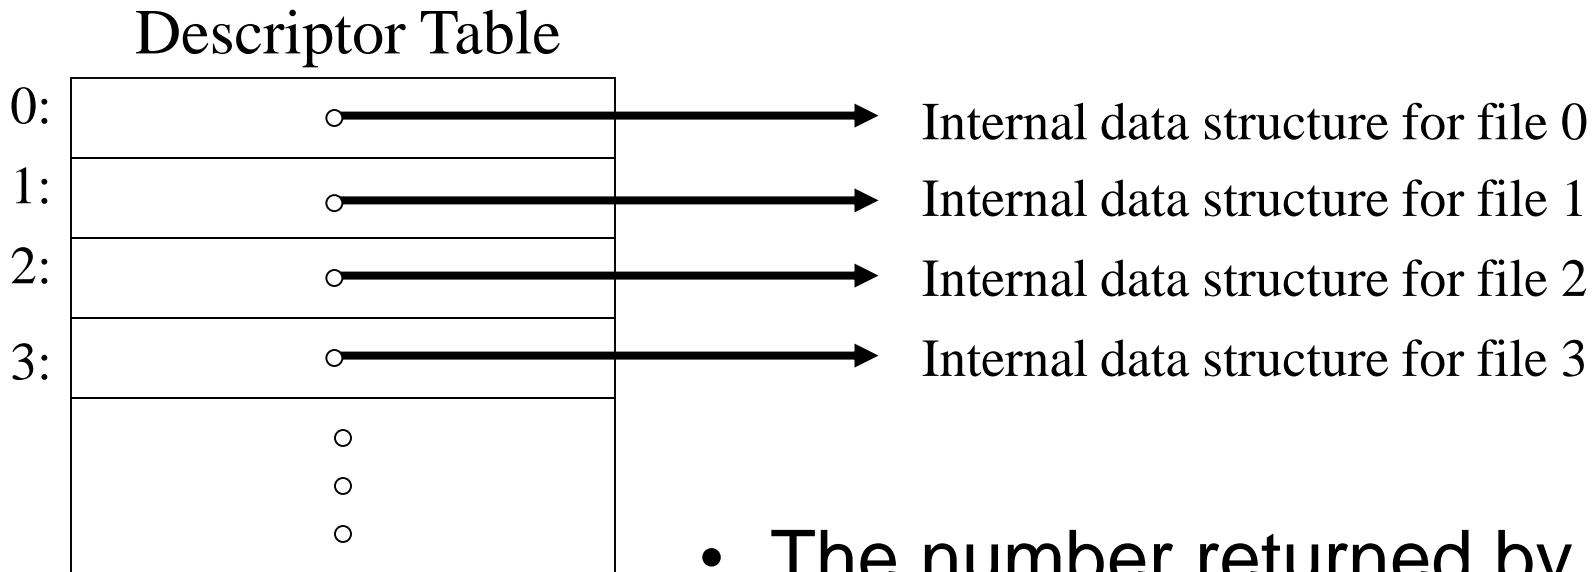
- UNIX Network **P**rogramming **I**nterfaces
 - X/Open Transport Interface (XTI)
 - Berkeley Sockets

Network Sockets



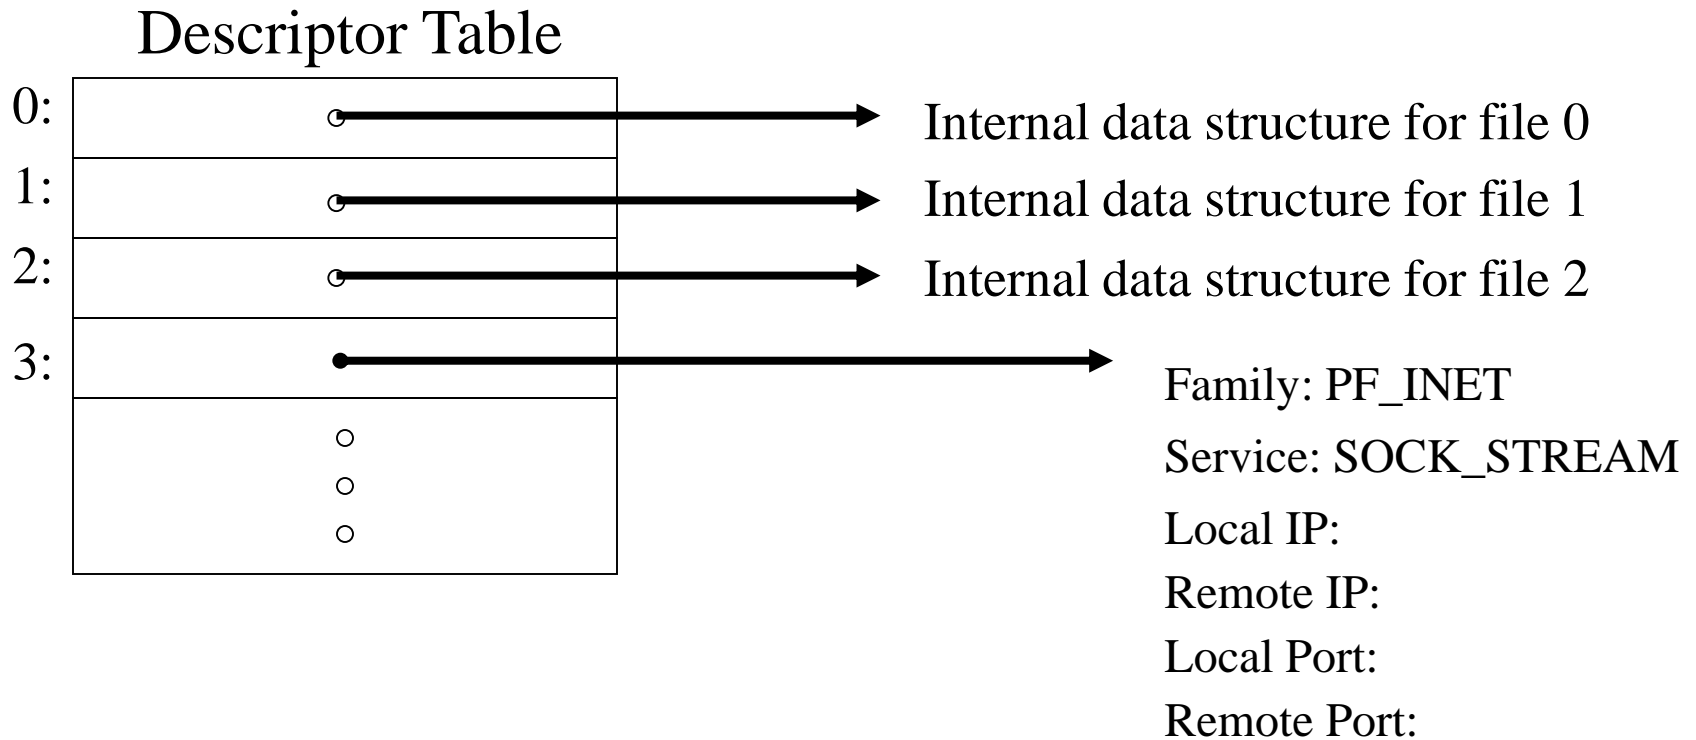
- **Berkeley Socket API**
 - A "socket" is the endpoint of a communication link between two processes
 - The socket API treats network connections like files as much as possible
 - Developed in the early 1980s for BSD Unix under a grant from DARPA

File Descriptor Table



- The number returned by `open ()` is an index into an array of pointers to internal OS data structures

File Descriptor Table



- Descriptor table after call to `socket()` command

Client Sockets & TCP/IP

- Client initializing connection with server
 - Server must already be running and expecting connections
- Since the client is using the TCP protocol, it needs to know the IP address and port of the server beforehand

Ports

- Many different processes can be running on one computer
- However, an IP address only identifies the machine, not the process
- Ports are used to reach a specific process on a machine
- Each process listens on a unique port - similar to a PO Box (ports are part of TCP)
- So a complete address is a IP address combined with a port number

Socket Documentation

- Most socket related man pages are in the "3n" section
 - `man -s 3n socket`
 - `man -k socket`
- All the info you need to use the network library is scattered across different man pages
 - Hard to use man pages for overall information
 - Best to have a book, and/or sample code available

Client Sockets

- Process:
 - 1. Create the socket endpoint - `socket()`
 - 2. Connect the socket to the server - `connect()`
 - 3. Use `read()` and `write()`, or `send()` and `recv()` to transfer data to and from the socket

Creating the Socket

```
int socket(int domain, int type, int protocol);
```

For IP, use AF_INET

Returns file
descriptor
or -1

For TCP, use SOCK_STREAM
For UDP, use SOCK_DGRAM

For IP, use 0

```
if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) == -1)
{
    perror("socket");
    exit(1);
}
```

Connecting the Socket

`connect()` connects the socket to a remote address

```
int connect(int sockfd,  
            struct sockaddr *address,  
            size_t address_size);
```

1. Sock file descriptor

2. Network address to connect to

3. Size of structure specified in 2.

Returns 0 on success or -1 on failure

```
if (connect(sockfd, (struct sockaddr *) &server, sizeof(server)) == -1)  
{  
    perror("connect");  
    exit(1);  
}
```

Filling the address struct

IP Number Address

- Client connecting to server

```
struct sockaddr_in server;
```

```
server.sin_family = AF_INET;
```

```
server.sin_port = htons(7000);
```

```
server.sin_addr.s_addr = inet_addr("192.168.1.1");
```

The IP address field

- Suppose you have declared

```
struct sockaddr_in server;
```

- If you acquire the address in the form of

```
in_addr_t
```

then you assign directly:

```
server.sin_addr.s_addr = inet_addr("192.168.1.1");
```

The IP address field

- If you have anything else, such as

```
int ipaddress;
```

- then you should be careful that the data doesn't get changed during type conversion
- `memcpy` is one solution:

```
memcpy(&server.sin_addr, &ipaddress, sizeof(ipaddress));
```

Setting up an address

- **Client connecting to server:**

```
struct hostent *server_ip_address;  
server_ip_address = gethostbyname("www.engr.orst.edu");  
  
if (server_ip_address == NULL)  
{  
    fprintf(stderr, "could not resolve server host name\n");  
    exit(1);  
}  
  
server.sin_family = AF_INET;  
server.sin_port = htons(80);  
  
memcpy(&server.sin_addr, server_ip_address->h_addr,  
       server_ip_address->h_length);
```

Setting up an address

- Server accepting connections:

```
struct sockaddr_in server;
```

```
server.sin_family = AF_INET;
```

```
server.sin_port = htons(7000);
```

```
server.sin_addr.s_addr = INADDR_ANY;
```


Sending data

1. Socket file descriptor

2. Pointer to data that should be sent

```
ssize_t send(int sockfd, void *message,  
             size_t message_size, int flags);
```

Returns number of bytes sent, or 0

3. Number of bytes to send, starting at address in 2

4. Configuration flags

```
char request[1024];  
r = send(sockfd, request, 1024, 0);  
if (r < 1024)  
    {} // handle possible error
```

Send

- Send will block until all the data has been sent, or the connection goes away
- Remember that internet connections fail all the time
 - Client intentionally disconnects (STOP button in a web browser)
 - Network partitions
 - Network failure
 - etc.

Receiving data

1. Sock file descriptor

2. Pointer to
buffer for
storing data

```
ssize_t recv(int sockfd, void *buffer,  
             size_t buffer_size, int flags);
```

Returns number
of bytes sent, or 0

3. Max number of bytes to
retrieve

4. Configuration
flags

```
char buffer[1024];  
r = recv(sockfd, buffer, 1024, 0);  
if (r < 1024)  
{  
    // error if r == -1  
    // if 0 < r < 1024, may be more data  
    // if r == 0, end of data  
}
```

Receiving data

- Data may arrive in odd size bundles
- `recv()` or `read()` will return exactly the amount of data that has already arrived
- more data may be coming as long as the return value is greater than 0
- `recv()` and `read()` will block if the connection is open but no data is available

Demo Client/Server

```
% gcc -o client client.c
```

```
% gcc -o server server.c
```

```
% ./server 51717 &
```

```
% ./client localhost 51717
```

Please enter the message: test

Here is the message: test

I got your message

```
[1]+ Done
```

```
./server 51717
```